

# Example: Mosaic in Action

## north carolina (simple + county)

Matt Mohn

2025-12-16

```
## Load necessary packages
library(sf)
library(tidyverse)
library(igraph)
library(sf)
library(readr)
library(ggplot2)
library(dplyr)
library(shadowtext)
library(scales)

## Load mosaic files
source("timing_functions.R")
source("scoring_functions.R")
source("tree_functions.R")
source("partition_functions.R")
source("recom_functions.R")
source("graph_functions.R")
source("output_functions.R")
source("chain_runner.R")
source("cleanup_globals.R")
source("graphics.R")
```

## Example 1: Simple Annealing

Let's setup and run the chain.

We'll use a generalized North Carolina VTD shapefile (generalization simplifies the shapefile - this doesn't affect the annealing in any way but makes plotting our results faster). It's saved here as `North_Carolina_Simplified.shp`

By default, Mosaic will run for 1000 iterations with a `pdev_tolerance` (acceptable population deviation) of  $\pm 5\%$ , and applies a cut edge weight of 1 (`weight_cut_edges = 1`). It will produce 5 districts by default, but we want North Carolina to have the 14 congressional districts it has in real life, so we set `num_districts` to 14 here.

```
SHAPEFILE <- "shapefiles/North_Carolina_Simplified.shp"

results <- run_chain(
  shapefile_path = SHAPEFILE,
  num_districts = 14,
  seed=123456
)

## Loading shapefile: shapefiles/North_Carolina_Simplified.shp
## Loading cached graph from: cache/cache_North_Carolina_Simplified.rds
## Loaded 2666 precincts
## Total population: 10,439,388
## Ideal per district: 745,671
## Population tolerance: 5.0%
##
## County data found: 100 unique counties
##
## Creating initial partition...
##
## GUIDED cooling: T0=167.6 to T=1.0 at iteration 900 (90%) at rate=0.994326
## Optimization enabled: Initial score=838.0, Temp=167.6
##
## Running 1000 ReCom steps...
##
## ITER      SCORE     TEMP    CUTS   CNTY   BUNK    MM DIFF   EF GAP    D      R      C      % ACC
## 100       792      94.9    792    --     -      --      --      --      --      --      --      92.0
## 200       704      53.7    704    --     -      --      --      --      --      --      --      92.0
## 300       708      30.4    708    --     -      --      --      --      --      --      --      90.0
## 400       741      17.2    741    --     -      --      --      --      --      --      --      84.0
##
## ITER      SCORE     TEMP    CUTS   CNTY   BUNK    MM DIFF   EF GAP    D      R      C      % ACC
## 500       631      9.7     631    --     -      --      --      --      --      --      --      66.0
## 600       683      5.5     683    --     -      --      --      --      --      --      --      46.0
## 700       588      3.1     588    --     -      --      --      --      --      --      --      36.0
## 800       561      1.8     561    --     -      --      --      --      --      --      --      22.0
## 900       555      1.0     555    --     -      --      --      --      --      --      --      10.0
##
## ITER      SCORE     TEMP    CUTS   CNTY   BUNK    MM DIFF   EF GAP    D      R      C      % ACC
## 1000      540      0.6     540    --     -      --      --      --      --      --      --      16.0
##
## Completed in 13.3s (13.3ms/step, 75.4 iter/s)
## Optimization stats: 563 accepts, 437 rejects (56.3% accept rate)
## Score: 838.0 to 540.0 (D=-298.0)
##
## =====
## FINAL ANALYSIS
```

```

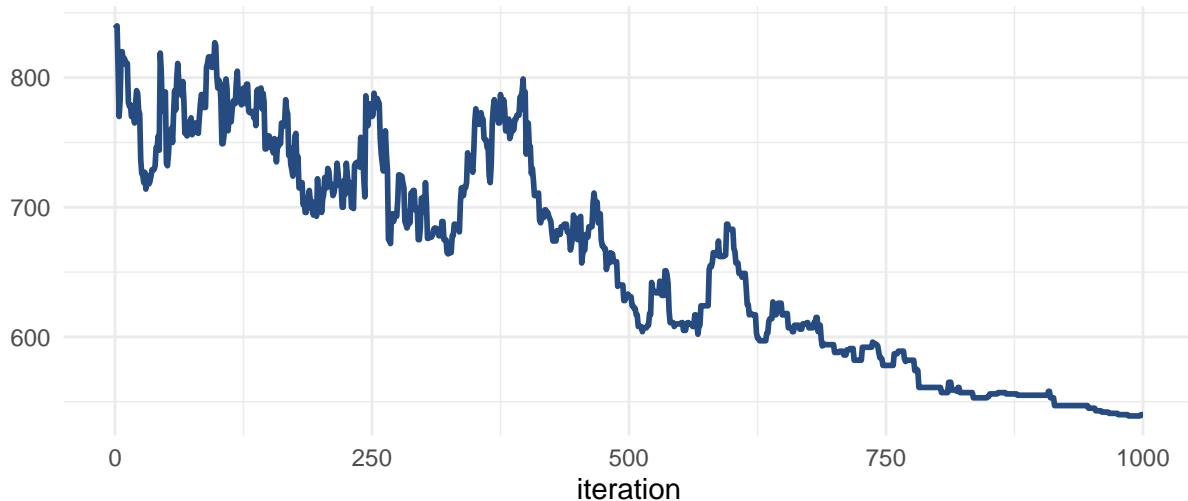
## =====
## =====
## =====
## RESULTS SAVED
## =====
## Final assignment: output/final_assignment_12182025_0033.csv
## Metrics: output/metrics_12182025_0033.csv
## Snapshots: output/snapshot_assignments_12182025_0033.csv
## =====

```

The results of the annealing are encouraging: the program initialized with 792 cuts (which isn't bad itself) and then fell to 540 over the course of the run.

The score perfectly matches the number of cuts because that's the only weight that we used.

Cut edges



*Knowing* the results is only a part of the fun. Let's see what the map actually looks like.

Mosaic includes the `mosaic_plot` file, which can create nice-looking maps on the fly.

```

# By default, mosaic_plot will call from the most recently
# saved final_assignments file.

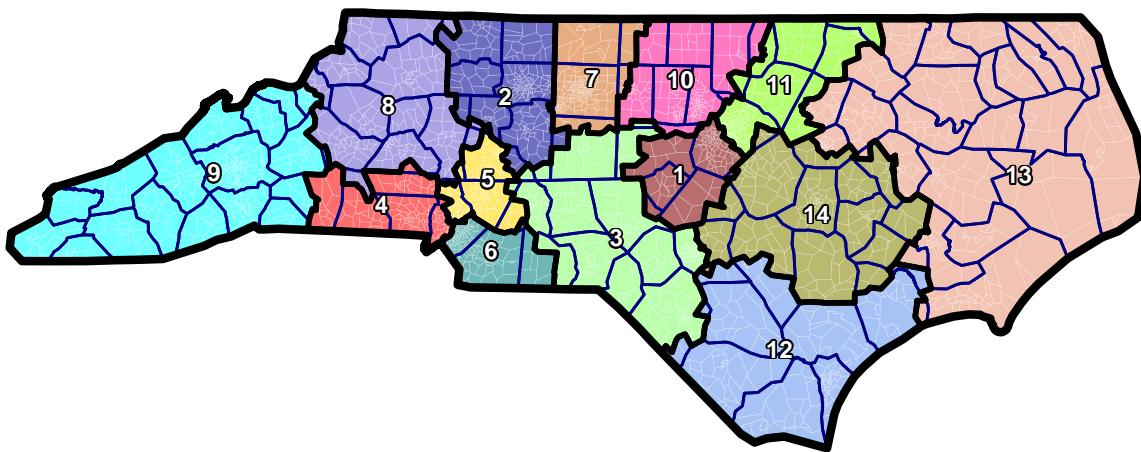
p <- mosaic_plot(shapefile_path = SHAPEFILE,
                  type="simple",
                  title="North Carolina: Example 1",
                  subtitle="Congressional districts, no weights",
                  border_outline = TRUE,
                  district_outline = TRUE,
                  county_outline = TRUE,
                  number_labels = TRUE,

```

```
precinct_outline = TRUE)  
p
```

## North Carolina: Example 1

Congressional districts, no weights



It looks great! All of the districts are fairly compact and there aren't many weird shapes. One issue jumps out though - we're splitting counties all over the place (look at Charlotte or Raleigh!) That makes sense- we didn't tell Mosaic to try in the first place. For that, let's move to example 2.

## Example 2: Preserving Counties

To better preserve counties, we take the same `run_chain()` call from earlier- but make two critical changes.

- First, we apply `county_bias`. The recombination algorithm usually randomly selects connections in the spanning graph to cut. However, we can induce it to draw cuts *between* counties, instead of *within* them, by increasing the sampling probability of edges that are also county boundaries. By setting `county_bias` to 10, we make it 10 times as likely that a county boundary is selected compared to a typical connection.
- Second, although `county_bias` will naturally lead to maps with fewer county splits, we also want the algorithm to try and value low-splitting as a general practice. We set `weight_county_splits` to 15. There are no units involved, but by default this means that 1 county split will be worth as much as 15 additional cut edges in annealing.

```
SHAPEFILE <- "shapefiles/North_Carolina_Simplified.shp"

results <- run_chain(
  shapefile_path = SHAPEFILE,
  num_districts = 14,
  seed=12345,
  county_bias=10,
  weight_county_splits=15,
)

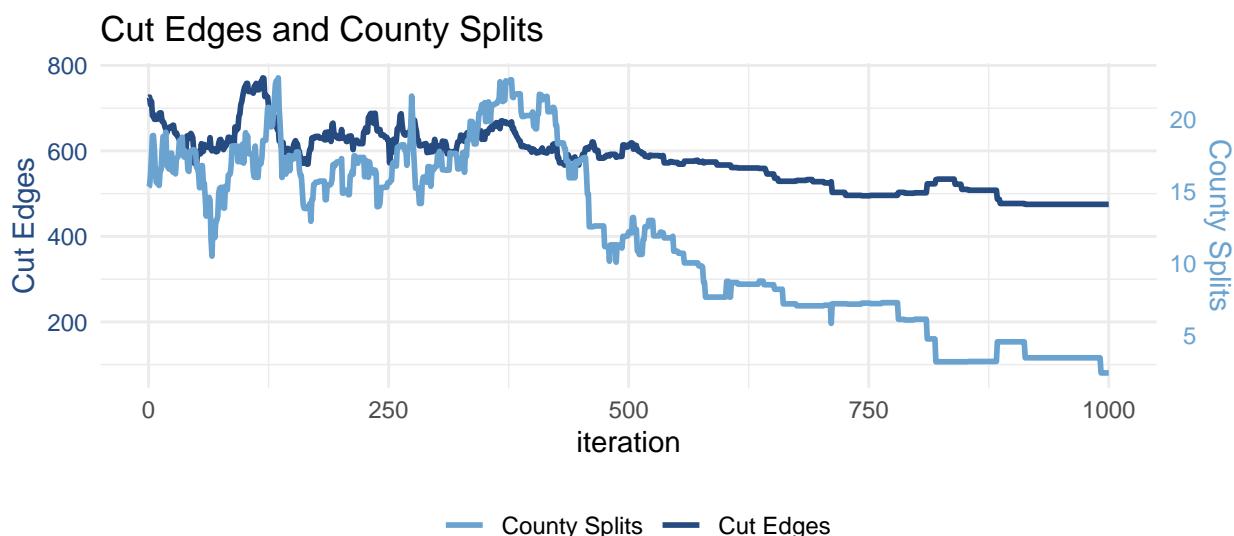
## Loading shapefile: shapefiles/North_Carolina_Simplified.shp
## Loading cached graph from: cache/cache_North_Carolina_Simplified.rds
## Loaded 2666 precincts
## Total population: 10,439,388
## Ideal per district: 745,671
## Population tolerance: 5.0%
##
## County data found: 100 unique counties
##
## Creating initial partition...
## GUIDED cooling: T0=191.0 to T=1.0 at iteration 900 (90%) at rate=0.994181
## Optimization enabled: Initial score=954.9, Temp=191.0
## County-aware mode: bias factor = 10.0
##
## Running 1000 ReCom steps...
##
## ITER      SCORE     TEMP    CUTS   CNTY   BUNK    MM DIFF   EF GAP    D      R      C      % ACC
## 100       1021     106.5   749    18.10   -      --      --      --      --      --      --      96.0
## 200       891      59.4    634    17.14   -      --      --      --      --      --      --      80.0
## 300       865      33.2    614    16.76   -      --      --      --      --      --      --      86.0
## 400       913      18.5    605    20.55   -      --      --      --      --      --      --      76.0
##
## ITER      SCORE     TEMP    CUTS   CNTY   BUNK    MM DIFF   EF GAP    D      R      C      % ACC
## 500       797      10.3    614    12.21   -      --      --      --      --      --      --      50.0
## 600       682      5.8     567    7.67    -      --      --      --      --      --      --      28.0
## 700       634      3.2     528    7.07    -      --      --      --      --      --      --      30.0
## 800       594      1.8     502    6.13    -      --      --      --      --      --      --      18.0
## 900       545      1.0     477    4.56    -      --      --      --      --      --      --      14.0
##
## ITER      SCORE     TEMP    CUTS   CNTY   BUNK    MM DIFF   EF GAP    D      R      C      % ACC
## 1000      511      0.6     475    2.40    -      --      --      --      --      --      --      20.0
##
## Completed in 14.4s (14.4ms/step, 69.7 iter/s)
```

```

## Optimization stats: 516 accepts, 484 rejects (51.6% accept rate)
## Score: 954.9 to 511.0 (D=-443.9)
##
## =====
## FINAL ANALYSIS
## =====
##
## County Splits Score: 2.40
## =====
##
##
## =====
## RESULTS SAVED
## =====
## Final assignment: output/final_assignment_12182025_0034.csv
## Metrics: output/metrics_12182025_0034.csv
## Snapshots: output/snapshot_assignments_12182025_0034.csv
## =====

```

The initial score (954) reflects the number of cut edges (725) plus 10 times the number of county splits (15.3). After 1000 iterations, the score has fallen to 511, with 475 cut edges and just 2.40 county splits.



For due diligence's sake, let's plot the map just to be sure.

```

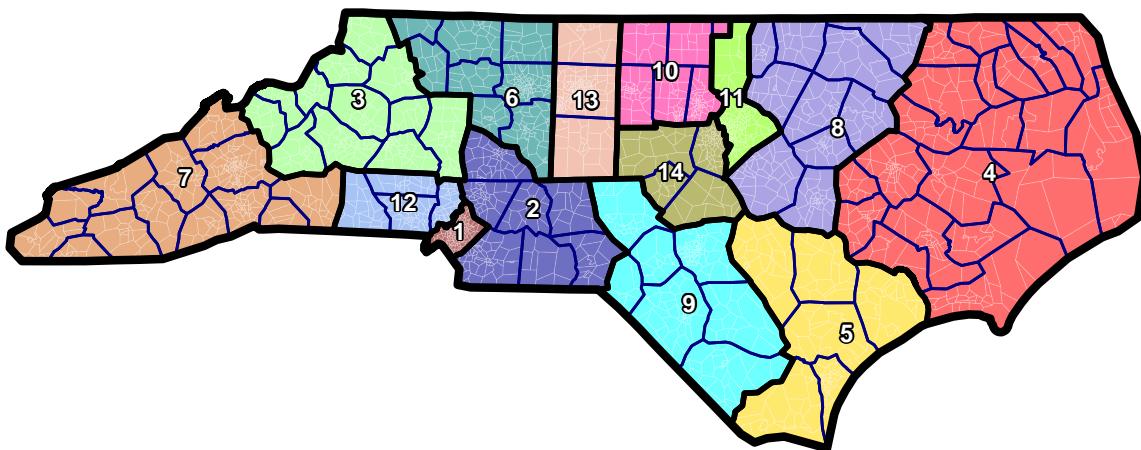
p <- mosaic_plot(shapefile_path = SHAPEFILE,
                  type="simple",
                  title="North Carolina: Example 2",
                  subtitle="Congressional districts with county weights",
                  border_outline = TRUE,
                  district_outline = TRUE,
                  county_outline = TRUE,
                  number_labels = TRUE,

```

```
precinct_outline = TRUE)  
p
```

## North Carolina: Example 2

Congressional districts with county weights



Looks good!

There are different ways of measuring county splits. Four counties are split here (Granville, Mecklenburg, Wake, and Wilkes) but because Meck and Wake are larger than a district each, we subtract them from the total, taking us down to 2.

The remaining 0.4 comes from a measure of fragmentation- this penalizes splits that are especially egregious.

Now that we know how to draw geographically nice-looking maps, Example 3 will explore using partisan data.

### Example 3: Republican Gerrymander

For Example 3, let's attempt to draw a Republican gerrymander- which, not coincidentally, is what the North Carolina General Assembly has been doing for the past two decades anyway. Learn from the best, right?

Let's use the following settings:

- Decrease `county_bias` to 2. First off, we care less about county splits this time around. Second, when we're looking to pull off partisan objectives, we don't want to constrain the set of possible combinations available to us.
- Decrease `weight_county_splits` to 5- again, nice to have, but not something mission-critical.
- Increase `num_steps` to 3000 up from the default 1000, to give us more room to maneuver.
- Set `target_efficiency_gap` to 0.5, meaning we aim for Republican votes to be 40% more efficient than Democratic ones. We weight this a 100 using `weight_efficiency_gap`.
- Set `target_mean_median` to 0.15, meaning we want the median district in North Carolina to be 15% more Republican by share of the vote than the state as a whole. If the state votes for a Democrat, we still want a majority to be won by Republicans. We weight this a 50 using `weight_mean_median`.
- Finally, set `target_dem_seats` to 2. This is pretty self-explanatory! We could aim for 0, but that would be very difficult and we don't want the score to be too meaningless. We weight this even further, to 250, using `weight_dem_seats`.

```
SHAPEFILE <- "shapefiles/North_Carolina_Simplified.shp"

results <- run_chain(
  shapefile_path = SHAPEFILE,
  num_districts = 14,
  seed=123456,
  county_bias=2,
  weight_county_splits=5,
  target_efficiency_gap=0.5,
  weight_efficiency_gap=100,
  target_mean_median=0.15,
  weight_mean_median=50,
  num_steps=3000,
  target_dem_seats=2,
  weight_dem_seats=250
)

## Loading shapefile: shapefiles/North_Carolina_Simplified.shp
## Loading cached graph from: cache/cache_North_Carolina_Simplified.rds
## Loaded 2666 precincts
## Total population: 10,439,388
## Ideal per district: 745,671
## Population tolerance: 5.0%
##
## County data found: 100 unique counties
##
## Creating initial partition...
## GUIDED cooling: T0=978.9 to T=1.0 at iteration 2700 (90%) at rate=0.997453
## Optimization enabled: Initial score=4894.7, Temp=978.9
## County-aware mode: bias factor = 2.0
##
## Running 3000 ReCom steps...
##
## ITER      SCORE      TEMP      CUTS      CNTY      BUNK      MM DIFF      EG ADJ      D      R      C      % ACC
## 100      4830      758.6      778      57.52      -      +1.9      +12.9      5.0      9.0      3.3      88.0
## 200      2982      587.8      724      53.18      -      +2.8      +20.9      4.1      9.9      2.2      90.0
## 300      4344      455.5      783      59.56      -      +2.2      +13.7      4.7      9.3      3.1      80.0
```

```

## 400      4527    352.9    775   64.84   -       +2.0     +14.6    4.9    9.1    3.7    78.0
##
## ITER    SCORE    TEMP     CUTS   CNTY    BUNK    MM DIFF   EG ADJ    D      R      C      % ACC
## 500      3797    273.5    766   61.73   -       +2.1     +16.9    4.5    9.5    2.9    86.0
## 600      3139    211.9    719   58.95   -       +3.3     +20.0    4.1    9.9    2.3    84.0
## 700      2982    164.2    632   52.99   -       +3.6     +19.9    4.1    9.9    2.2    88.0
## 800      3006    127.2    736   52.22   -       +3.4     +19.8    4.0    10.0   2.1    82.0
## 900      3015    98.6     801   57.14   -       +3.1     +21.1    4.0    10.0   2.1    70.0
##
## ITER    SCORE    TEMP     CUTS   CNTY    BUNK    MM DIFF   EG ADJ    D      R      C      % ACC
## 1000     2713    76.4     737   47.47   -       +3.1     +22.0    3.9    10.1   1.8    68.0
## 1100     2864    59.2     681   56.51   -       +2.5     +21.2    4.0    10.0   2.0    70.0
## 1200     2811    45.9     684   52.59   -       +3.1     +22.4    4.0    10.0   2.1    66.0
## 1300     2653    35.5     716   48.09   -       +4.0     +22.6    3.9    10.1   1.8    48.0
## 1400     2519    27.5     682   50.08   -       +3.9     +23.1    3.8    10.2   1.6    54.0
##
## ITER    SCORE    TEMP     CUTS   CNTY    BUNK    MM DIFF   EG ADJ    D      R      C      % ACC
## 1500     2411    21.3     675   46.13   -       +3.8     +24.1    3.8    10.2   1.5    50.0
## 1600     2388    16.5     656   46.48   -       +3.6     +23.8    3.7    10.3   1.5    48.0
## 1700     2450    12.8     639   41.99   -       +3.2     +23.3    3.8    10.2   1.6    44.0
## 1800     2174    9.9      674   38.77   -       +5.1     +24.5    3.6    10.4   1.1    38.0
## 1900     2227    7.7      708   42.09   -       +4.7     +24.3    3.5    10.5   1.1    24.0
##
## ITER    SCORE    TEMP     CUTS   CNTY    BUNK    MM DIFF   EG ADJ    D      R      C      % ACC
## 2000     2168    6.0      655   42.56   -       +4.4     +24.5    3.5    10.5   1.1    22.0
## 2100     2156    4.6      694   42.60   -       +4.6     +24.8    3.5    10.5   1.0    14.0
## 2200     2110    3.6      656   40.70   -       +4.7     +25.1    3.5    10.5   1.0    12.0
## 2300     2105    2.8      658   40.62   -       +4.7     +24.9    3.5    10.5   1.0    16.0
## 2400     2069    2.1      636   35.20   -       +5.0     +24.7    3.5    10.5   1.0    8.0
##
## ITER    SCORE    TEMP     CUTS   CNTY    BUNK    MM DIFF   EG ADJ    D      R      C      % ACC
## 2500     2062    1.7      623   34.61   -       +4.7     +24.7    3.5    10.5   1.0    10.0
## 2600     2055    1.3      623   33.36   -       +4.8     +24.7    3.5    10.5   1.0    6.0
## 2700     2043    1.0      620   33.57   -       +5.0     +24.7    3.5    10.5   1.0    8.0
## 2800     2043    0.8      627   33.22   -       +5.0     +24.8    3.5    10.5   1.0    2.0
## 2900     2045    0.6      631   33.23   -       +5.0     +24.8    3.5    10.5   1.0    6.0
##
## ITER    SCORE    TEMP     CUTS   CNTY    BUNK    MM DIFF   EG ADJ    D      R      C      % ACC
## 3000     2045    0.5      631   33.23   -       +5.0     +24.8    3.5    10.5   1.0    8.0
##
## Completed in 49.0s (16.3ms/step, 61.3 iter/s)
## Optimization stats: 1361 accepts, 1639 rejects (45.4% accept rate)
## Score: 4894.7 to 2044.6 (D=-2850.2)
##
## =====
## FINAL ANALYSIS
## =====
## District      DEM      REP      DEM%     REP%     Margin    P(DEM)  Winner
## -----
## D1            165791   224162   42.5%    57.5%    15.0%    3.6% REP
## D2            279827   161550   63.4%    36.6%    26.8%    99.7% DEM
## D3            160533   198429   44.7%    55.3%    10.6%    9.0% REP
## D4            175933   228023   43.6%    56.4%    12.9%    5.6% REP
## D5            171118   224416   43.3%    56.7%    13.5%    4.9% REP
## D6            168377   222809   43.0%    57.0%    13.9%    4.5% REP
## D7            170273   227545   42.8%    57.2%    14.4%    4.1% REP
## D8            160673   241431   40.0%    60.0%    20.1%    1.2% REP
## D9            190138   236741   44.5%    55.5%    10.9%    8.3% REP
## D10           167491   213946   43.9%    56.1%    12.2%    6.4% REP

```

```

## D11          323387    114775    73.8%    26.2%    47.6%    100.0% DEM
## D12          250579     87110    74.2%    25.8%    48.4%    100.0% DEM
## D13          172067    243616    41.4%    58.6%    17.2%     2.2% REP
## D14          157422    272388    36.6%    63.4%    26.7%     0.3% REP
## -----
## TOTAL        2713609   2896941    48.4%    51.6%
##
## Mean Dem Share: 48.4%
## Median Dem Share: 43.4%
## MM: R+5.0% (tgt R+15.0%) | MM Score: 1
## EG: R+24.8% (tgt R+50.0%) | EG Score: 6
## Expected Seats: D 3.5 | R 10.5 (tgt D=2.0) | Seats Score: 2
##
## County Splits Score: 33.23
## =====
##
##
## =====
## RESULTS SAVED
## =====
## Final assignment: output/final_assignment_12182025_0035.csv
## Metrics: output/metrics_12182025_0035.csv
## Snapshots: output/snapshot_assignments_12182025_0035.csv
## =====

```

The results are encouraging. MM DIFF is R+5.0 (not R+15, our target, but pretty good for 3000 runs), and EG ADJ (efficiency gap adjusted over a few different swing scenarios) hit R+25.

The model expects Republicans to win 10.5 seats over the long-run average, with Democrats winning only 3.5 - but the seat table makes it pretty clear that in most years this is a 11-3 map.

Once again, we graph using `mosaic_plot` - this time, however, we can set `type` to 'partisan' and look at the districts' final partisanship. This only works if we ran the annealing using partisan data- otherwise the annealing will not store this information.

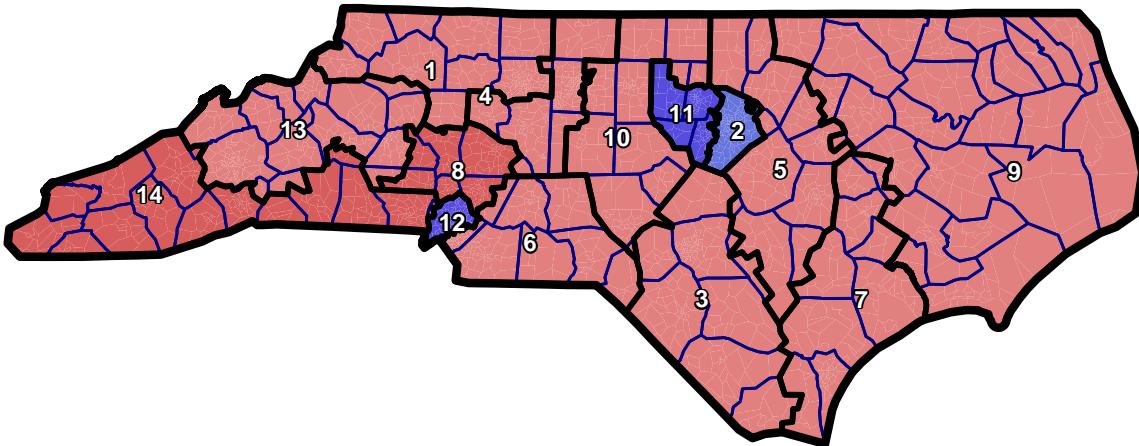
```

p <- mosaic_plot(shapefile_path = SHAPEFILE,
                  type="partisan",
                  title="North Carolina: Example 3",
                  subtitle="Republican congressional gerrymander",
                  border_outline = TRUE,
                  district_outline = TRUE,
                  county_outline = TRUE,
                  number_labels = TRUE,
                  precinct_outline = FALSE)
p

```

## North Carolina: Example 3

Republican congressional gerrymander



Anyone familiar with North Carolina gerrymandering history will immediately recognize this kind of map.

- The Triad is cracked; Greensboro and Winston-Salem are put into separate districts. District 1 (the Winston-to-TN district) looks like Virginia Foxx's NC-5 from 2001-2021. Greensboro itself is split perfectly in half.
- Charlotte, Raleigh, Chapel Hill, and Durham are all packed into deep blue districts.
- The Black Belt is split. Goldsboro, Kinston, and Rocky Mount all end up in three different districts.
- Two 'claws' form around the Triangle (Districts 5 and 10) instead of a district forming along the VA border.
- Fayetteville's district reaches south (towards South Carolina). Any other direction would be more purple.
- The purple-ish suburbs around Charlotte end up in three different districts.

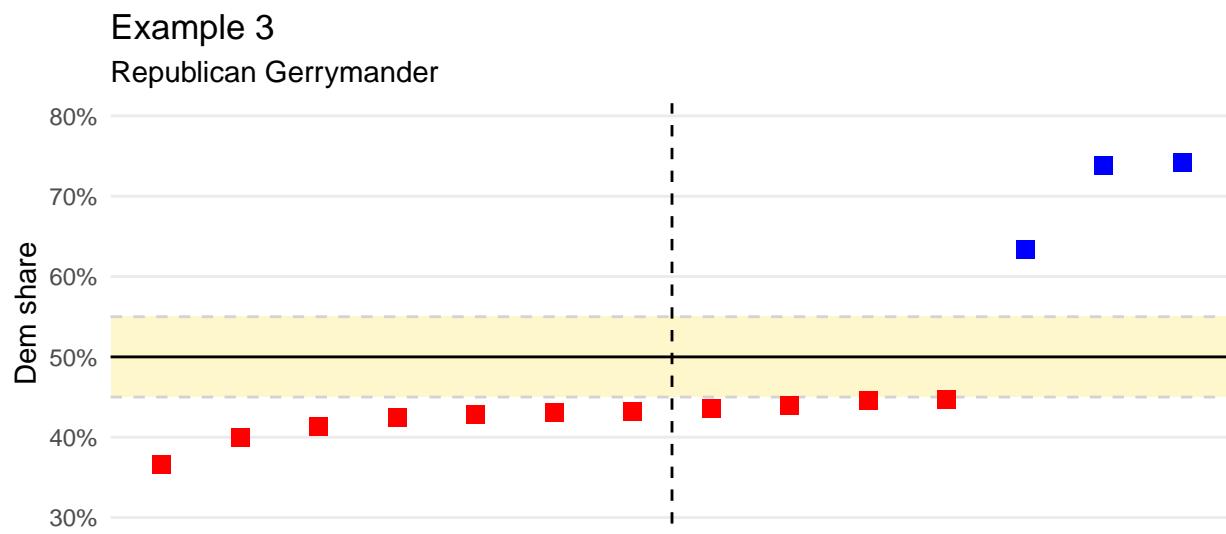
All of these choices (major and minor) contribute to this kind of partisan lean, but the algorithm figured out unknowingly what took Republican mapmakers a while to master.

Let's check out the partisan skew using `mosaic_partisan_plot`

```
p <- mosaic_partisan_plot(title="Example 3",
                           subtitle="Republican Gerrymander")

## Loading data...
## Using most recent metrics CSV: output/metrics_12182025_0035.csv
## Loading metrics for partisan plot...
## Rendering partisan graph plot...
```

p



Gerrymandering in North Carolina is a bipartisan tradition, however, so for the reverse take, check out Example 4.

## Example 4: Democratic Gerrymander

For the Democratic attempt, we will basically reverse the settings from Example 3. The changes are:

- Set `target_efficiency_gap` to -0.5, meaning we aim for **Democratic** votes to be 50% more efficient than Republican ones.
- Set `target_mean_median` to -0.1, meaning we want the median district in North Carolina to be 15% less Republican by share of the vote than the state as a whole.
- Finally, set `target_dem_seats` to 10.

For concision, let's suppress the console output summaries with `verbose_console = FALSE`.

```
SHAPEFILE <- "shapefiles/North_Carolina_Simplified.shp"

results <- run_chain(
  shapefile_path = SHAPEFILE,
  num_districts = 14,
  seed=123456,
  county_bias=2,
  weight_county_splits=5,
  target_efficiency_gap=-0.3,
  weight_efficiency_gap=100,
  target_mean_median=-0.1,
  weight_mean_median=50,
  num_steps=3000,
  target_dem_seats=11,
  weight_dem_seats=250,
  verbose_console=FALSE
)

## Loading shapefile: shapefiles/North_Carolina_Simplified.shp
## Loading cached graph from: cache/cache_North_Carolina_Simplified.rds
## Loaded 2666 precincts
## Total population: 10,439,388
## Ideal per district: 745,671
## Population tolerance: 5.0%
##
## County data found: 100 unique counties
##
## Creating initial partition...
## GUIDED cooling: T0=2356.7 to T=1.0 at iteration 2700 (90%) at rate=0.997128
## Optimization enabled: Initial score=11783.5, Temp=2356.7
## County-aware mode: bias factor = 2.0
##
## Running 3000 ReCom steps...
## Iter: 100/3000
## Iter: 200/3000
## Iter: 300/3000
## Iter: 400/3000
## Iter: 500/3000
## Iter: 600/3000
## Iter: 700/3000
## Iter: 800/3000
## Iter: 900/3000
## Iter: 1000/3000
## Iter: 1100/3000
## Iter: 1200/3000
## Iter: 1300/3000
## Iter: 1400/3000
```

```

## Iter: 1500/3000
## Iter: 1600/3000
## Iter: 1700/3000
## Iter: 1800/3000
## Iter: 1900/3000
## Iter: 2000/3000
## Iter: 2100/3000
## Iter: 2200/3000
## Iter: 2300/3000
## Iter: 2400/3000
## Iter: 2500/3000
## Iter: 2600/3000
## Iter: 2700/3000
## Iter: 2800/3000
## Iter: 2900/3000
## Iter: 3000/3000
##
##
## Completed in 49.2s (16.4ms/step, 61.0 iter/s)
## Optimization stats: 1318 accepts, 1682 rejects (43.9% accept rate)
## Score: 11783.5 to 3454.8 (D=-8328.6)
##
## =====
## FINAL ANALYSIS
## =====
## District      DEM      REP     DEM%    REP%   Margin   P(DEM) Winner
## -----
## D1            131802   290996  31.2%   68.8%  37.7%   0.0% REP
## D2            109458   281488  28.0%   72.0%  44.0%   0.0% REP
## D3            236940   176329  57.3%   42.7%  14.7%  96.2% DEM
## D4            193716   223180  46.5%   53.5%  7.1%   17.5% REP
## D5            219294   166322  56.9%   43.1%  13.7%  95.3% DEM
## D6            236493   172669  57.8%   42.2%  15.6%  96.9% DEM
## D7            187239   267543  41.2%   58.8%  17.7%   2.0% REP
## D8            244807   155924  61.1%   38.9%  22.2%  99.2% DEM
## D9            200467   148919  57.4%   42.6%  14.8%  96.2% DEM
## D10           210725   150713  58.3%   41.7%  16.6%  97.5% DEM
## D11           127459   292161  30.4%   69.6%  39.3%   0.0% REP
## D12           145212   256104  36.2%   63.8%  27.6%   0.2% REP
## D13           234138   160019  59.4%   40.6%  18.8%  98.4% DEM
## D14           235859   154574  60.4%   39.6%  20.8%  99.0% DEM
## -----
## TOTAL         2713609  2896941  48.4%   51.6%
##
## Mean Dem Share: 48.7%
## Median Dem Share: 57.1%
## MM: D+8.4% (tgt D+10.0%) | MM Score: 0
## EG: D+10.7% (tgt D+30.0%) | EG Score: 4
## Expected Seats: D 8.0 | R 6.0 (tgt D=11.0) | Seats Score: 9
##
## County Splits Score: 33.59
## =====
## 
## =====
## RESULTS SAVED
## =====
## Final assignment: output/final_assignment_12182025_0036.csv
## Metrics: output/metrics_12182025_0036.csv

```

```
## Snapshots: output/snapshot_assignments_12182025_0036.csv
## =====
```

The results look great, but in the opposite direction from earlier. MM DIFF is D+8.4, meaning the median district gave Kamala Harris 57.1% of the vote, while the average across the state was just 48.7%.

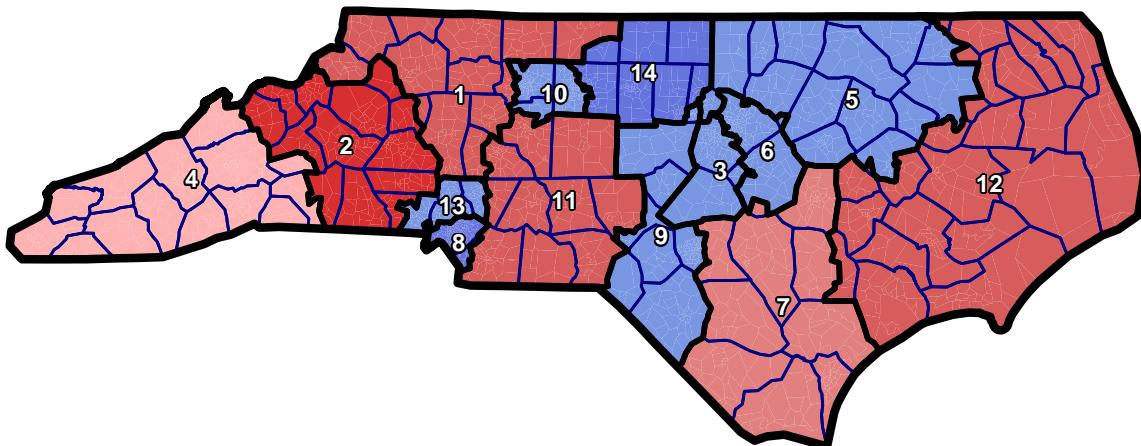
EG ADJ is D+10.7, and the model expects Democrats to win 8.0 seats and Republicans 6.0.

Again, we graph using `mosaic_plot`.

```
p <- mosaic_plot(shapefile_path = SHAPEFILE,
                  type="partisan",
                  title="North Carolina: Example 4",
                  subtitle="Democratic congressional gerrymander",
                  border_outline = TRUE,
                  district_outline = TRUE,
                  county_outline = TRUE,
                  number_labels = TRUE,
                  precinct_outline = FALSE)
p
```

## North Carolina: Example 4

### Democratic congressional gerrymander



Here, many of the opposite decisions from the Republican map have been taken.

- Charlotte is cracked in two- not packed. The resulting districts (12 and 6) both vote solidly blue.
- Similarly, the Triangle is cracked- not packed- and its blue votes are distributed into 5 different districts.

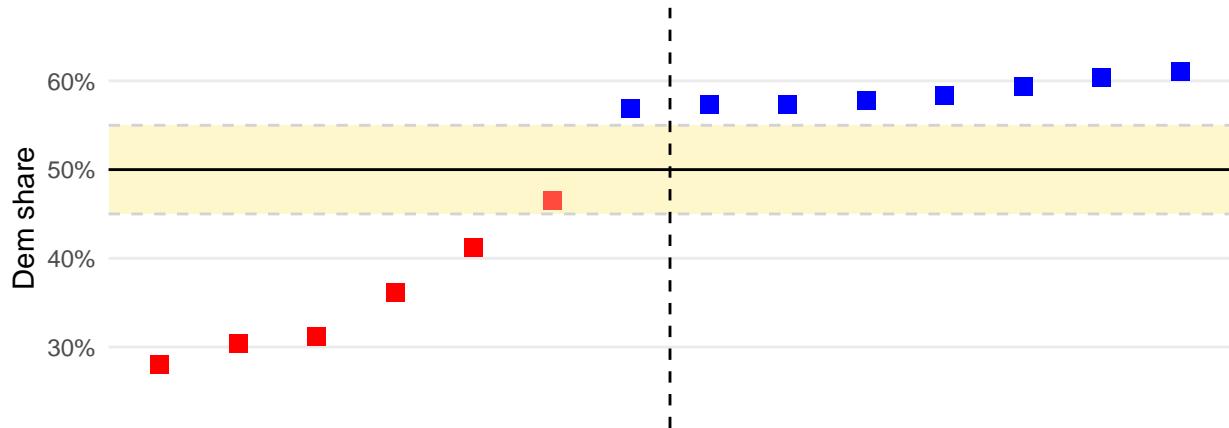
- The Piedmont Triad is mostly united in District 10 (with some of Greensboro's excess blue elsewhere). Fayetteville's district aims north to purple Raleigh suburbs, not east.

On the Tennessee border, Mosaic sees a chance at a competitive district, and while it can't draw a blue one, it follows the ghost of Heath Shuler to come up with one (District 4) that Trump only won by 7 points.

```
## Loading data...
## Using most recent metrics CSV: output/metrics_12182025_0036.csv
## Loading metrics for partisan plot...
## Rendering partisan graph plot...
```

#### Example 4

##### Democratic Gerrymander



## Example 5: A Fair State Senate Map

Let's deploy the partisan tools for a good cause- drawing a fair/competitive North Carolina State Senate map.

- Switch `target_mean_median` to 0 and `target_efficiency_gap` to 0. This is the same as favoring neither party.
- Introduce `weight_competitiveness` to 15. This metric tries to increase the share of districts that are competitive, where a 100% competitive district is defined as having a 50/50 chance of going to either party.
- Bolster our county tools, like `county_bias` and `weight_county_splits` - since we're after noble causes now.

```
SHAPEFILE <- "shapefiles/North_Carolina_Simplified.shp"

results <- run_chain(
  shapefile_path = SHAPEFILE,
  num_districts = 50,
  seed=123456,
  county_bias=5,
  weight_county_splits=10,
  target_mean_median=0,
  weight_mean_median=150,
  target_efficiency_gap=0,
  weight_efficiency_gap=100,
  weight_competitiveness=5,
  num_steps=5000,
  verbose_console = FALSE
)

## Loading shapefile: shapefiles/North_Carolina_Simplified.shp
## Loading cached graph from: cache/cache_North_Carolina_Simplified.rds
## Loaded 2666 precincts
## Total population: 10,439,388
## Ideal per district: 208,788
## Population tolerance: 5.0%
##
## County data found: 100 unique counties
##
## Creating initial partition...
## GUIDED cooling: T0=519.8 to T=1.0 at iteration 4500 (90%) at rate=0.998611
## Optimization enabled: Initial score=2599.0, Temp=519.8
## County-aware mode: bias factor = 5.0
##
## Running 5000 ReCom steps...
## Iter: 100/5000
## Iter: 200/5000
## Iter: 300/5000
## Iter: 400/5000
## Iter: 500/5000
## Iter: 600/5000
## Iter: 700/5000
## Iter: 800/5000
## Iter: 900/5000
## Iter: 1000/5000
## Iter: 1100/5000
## Iter: 1200/5000
```

```

## Iter: 1300/5000
## Iter: 1400/5000
## Iter: 1500/5000
## Iter: 1600/5000
## Iter: 1700/5000
## Iter: 1800/5000
## Iter: 1900/5000
## Iter: 2000/5000
## Iter: 2100/5000
## Iter: 2200/5000
## Iter: 2300/5000
## Iter: 2400/5000
## Iter: 2500/5000
## Iter: 2600/5000
## Iter: 2700/5000
## Iter: 2800/5000
## Iter: 2900/5000
## Iter: 3000/5000
## Iter: 3100/5000
## Iter: 3200/5000
## Iter: 3300/5000
## Iter: 3400/5000
## Iter: 3500/5000
## Iter: 3600/5000
## Iter: 3700/5000
## Iter: 3800/5000
## Iter: 3900/5000
## Iter: 4000/5000
## Iter: 4100/5000
## Iter: 4200/5000
## Iter: 4300/5000
## Iter: 4400/5000
## Iter: 4500/5000
## Iter: 4600/5000
## Iter: 4700/5000
## Iter: 4800/5000
## Iter: 4900/5000
## Iter: 5000/5000
##
##
## Completed in 116.3s (23.3ms/step, 43.0 iter/s)
## Optimization stats: 3522 accepts, 1478 rejects (70.4% accept rate)
## Score: 2599.0 to 1896.3 (D=-702.7)
##
## =====
## FINAL ANALYSIS
## =====
## District      DEM      REP      DEM%      REP%      Margin      P(DEM)  Winner
## -----
## D1          56550    47027    54.6%    45.4%     9.2%    88.3%  DEM
## D2          53863    93152    36.6%    63.4%    26.7%     0.3%  REP
## D3          43195    75551    36.4%    63.6%    27.2%     0.3%  REP
## D4          39169    74003    34.6%    65.4%    30.8%     0.1%  REP
## D5          52431    43802    54.5%    45.5%     9.0%    87.8%  DEM
## D6          53660    53195    50.2%    49.8%     0.4%    52.4%  DEM
## D7          61162    60527    50.3%    49.7%     0.5%    52.9%  DEM
## D8          66218    60327    52.3%    47.7%     4.7%    73.6%  DEM
## D9          51658    61485    45.7%    54.3%     8.7%    12.9%  REP
## D10         48476    72286    40.1%    59.9%    19.7%     1.3%  REP

```

## D11	73798	52967	58.2%	41.8%	16.4%	97.4% DEM
## D12	68050	38586	63.8%	36.2%	27.6%	99.8% DEM
## D13	57438	56645	50.3%	49.7%	0.7%	53.8% DEM
## D14	56269	53020	51.5%	48.5%	3.0%	65.8% DEM
## D15	75511	40220	65.2%	34.8%	30.5%	99.9% DEM
## D16	32604	78236	29.4%	70.6%	41.2%	0.0% REP
## D17	75725	41641	64.5%	35.5%	29.0%	99.8% DEM
## D18	35217	83591	29.6%	70.4%	40.7%	0.0% REP
## D19	93253	16854	84.7%	15.3%	69.4%	100.0% DEM
## D20	45663	52867	46.3%	53.7%	7.3%	16.7% REP
## D21	59593	68040	46.7%	53.3%	6.6%	18.9% REP
## D22	71095	28310	71.5%	28.5%	43.0%	100.0% DEM
## D23	76768	45691	62.7%	37.3%	25.4%	99.6% DEM
## D24	44111	83913	34.5%	65.5%	31.1%	0.1% REP
## D25	62079	30942	66.7%	33.3%	33.5%	99.9% DEM
## D26	41429	69566	37.3%	62.7%	25.3%	0.4% REP
## D27	32750	78065	29.6%	70.4%	40.9%	0.0% REP
## D28	62418	24942	71.4%	28.6%	42.9%	100.0% DEM
## D29	34411	73341	31.9%	68.1%	36.1%	0.0% REP
## D30	25684	54960	31.8%	68.2%	36.3%	0.0% REP
## D31	42241	41203	50.6%	49.4%	1.2%	56.8% DEM
## D32	75689	17378	81.3%	18.7%	62.7%	100.0% DEM
## D33	72150	55207	56.7%	43.3%	13.3%	94.9% DEM
## D34	34782	92695	27.3%	72.7%	45.4%	0.0% REP
## D35	30280	89453	25.3%	74.7%	49.4%	0.0% REP
## D36	59561	60780	49.5%	50.5%	1.0%	44.5% REP
## D37	71258	36456	66.2%	33.8%	32.3%	99.9% DEM
## D38	82511	44235	65.1%	34.9%	30.2%	99.9% DEM
## D39	76100	28712	72.6%	27.4%	45.2%	100.0% DEM
## D40	40844	67933	37.5%	62.5%	24.9%	0.4% REP
## D41	39401	78008	33.6%	66.4%	32.9%	0.1% REP
## D42	65865	64316	50.6%	49.4%	1.2%	56.5% DEM
## D43	42831	52120	45.1%	54.9%	9.8%	10.4% REP
## D44	32455	75241	30.1%	69.9%	39.7%	0.0% REP
## D45	39199	70148	35.8%	64.2%	28.3%	0.2% REP
## D46	74098	44908	62.3%	37.7%	24.5%	99.5% DEM
## D47	41580	80313	34.1%	65.9%	31.8%	0.1% REP
## D48	46911	71808	39.5%	60.5%	21.0%	1.0% REP
## D49	60227	64314	48.4%	51.6%	3.3%	32.7% REP
## D50	35378	47961	42.5%	57.5%	15.1%	3.5% REP
## -----						
## TOTAL	2713609	2896941	48.4%	51.6%		

## Mean Dem Share: 48.7%

## Median Dem Share: 48.9%

## MM: D+0.2% (tgt R+0.0%) | MM Score: 0

## EG: R+2.1% (tgt R+0.0%) | EG Score: 0

## Expected Seats: D 22.2 | R 27.8

##

## County Splits Score: 27.48

## =====

##

##

## =====

## RESULTS SAVED

## =====

## Final assignment: output/final\_assignment\_12182025\_0038.csv

## Metrics: output/metrics\_12182025\_0038.csv

## Snapshots: output/snapshot\_assignments\_12182025\_0038.csv

```
## =====
```

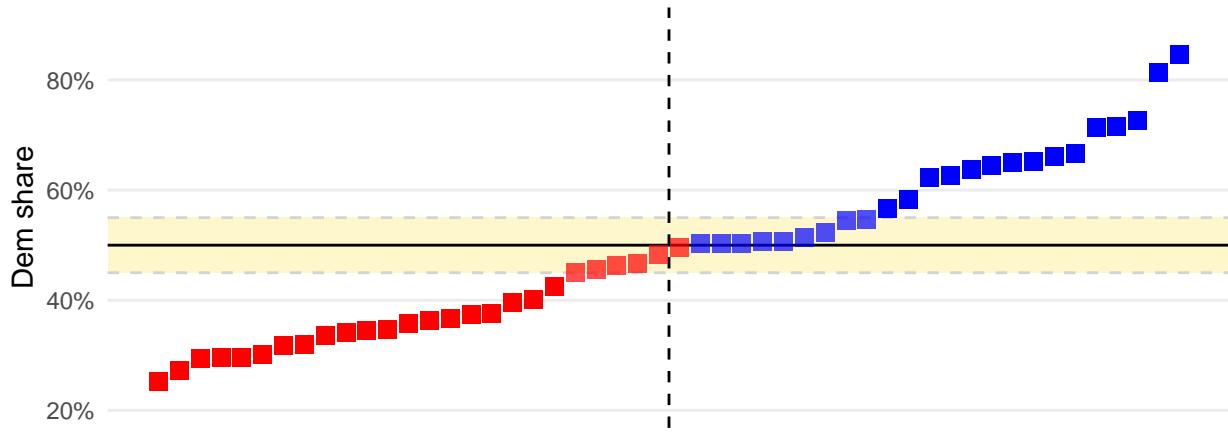
The results are compelling.

- The mean Democratic vote share (using Kamala Harris' 2024 result) is 48.7%, and the median is 48.9%. Trump wins a majority of districts, but if Harris had won, she likely would have too.
- The efficiency gap is R+2.1%, well within our tolerance.
- 9.3 competitive seats are forecast- almost 20% of the total map- creating a wide battleground.

```
## Loading data...
## Using most recent metrics CSV: output/metrics_12182025_0038.csv
## Loading metrics for partisan plot...
## Rendering partisan graph plot...
```

## Example 5

### Fair NC State Senate

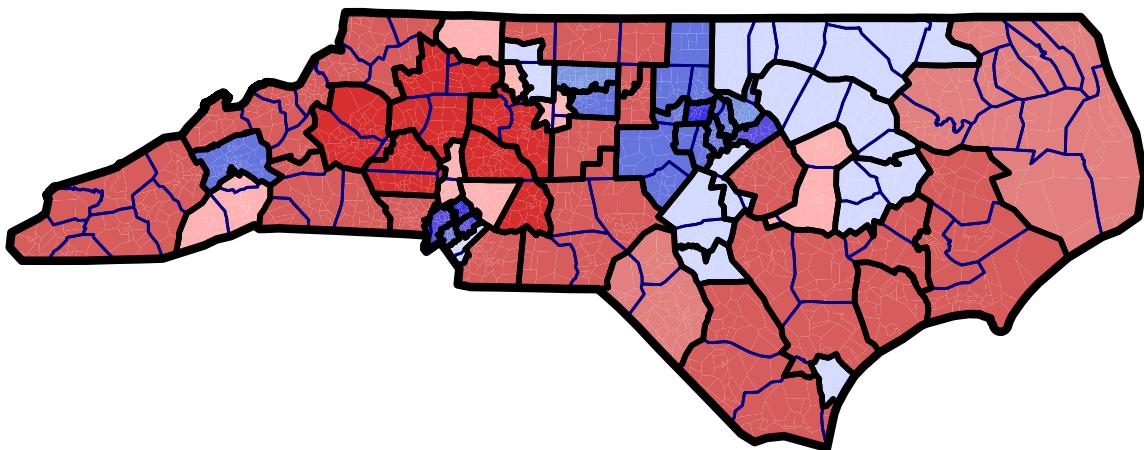


... and visualize! Note- in reality, this map would not pass the extremely strict splitting laws that North Carolina has in place for the NCGA. But it's fun to imagine.

```
p <- mosaic_plot(shapefile_path = SHAPEFILE,
                  type="partisan",
                  title="North Carolina: Example 5",
                  subtitle="Balanced and competitive NC Senate map",
                  border_outline = TRUE,
                  district_outline = TRUE,
                  county_outline = TRUE,
                  number_labels = FALSE,
                  precinct_outline = FALSE)
p
```

## North Carolina: Example 5

Balanced and competitive NC Senate map



## Example 6: Dealing with Incumbents

Let's imagine that tomorrow a court orders that North Carolina draw a new, fair congressional map. *However*, the court orders that the remedial map try and keep incumbents apart.

Mosaic can help here. The program can handle both pro-bunking and anti-bunking lists. Pro-bunking lists aim to keep precincts together (like communities of interest). Anti-bunking lists aim to keep precincts apart (like incumbent homes).

We establish the incumbent list here. These are the precinct IDs (GEOIDs) for each North Carolina congressperson (roughly - we don't really know where they live). The first two values are the weight and the exponent (here, we go with 1000 and 1 respectively).

```
NC_incumbent_list <- c(
  1000,                      # Weight: heavy penalty for splitting
  1,                          # Exponent: linear penalty
  "3707900BEAR", #Davis
  "37183001-27", #Ross
  "3714701504A", #Murphy
  "3713500000H", #Foushee
  "37011000002", #Foxx
  "37059000013", #McDowell
  "37129000W25", #Rouzer
  "37179000014", #Harris
  "37125000SSP", #Hudson
  "37035000013", #Harrigan
  "370890000FR", #Edwards
  "37119000011", #Adams
  "37183001-36", #Knott
  "3704500KM-N" #Moore
)
```

We pass the bunking list by adding a call to `create_bunking_lists()` to `run_chain()`. Multiple sub-lists can be added to either `anti_bunking` or `pro_bunking`.

```
## Loading shapefile: shapefiles/North_Carolina_Simplified.shp
## Loading cached graph from: cache/cache_North_Carolina_Simplified.rds
## Loaded 2666 precincts
## Total population: 10,439,388
## Ideal per district: 745,671
## Population tolerance: 5.0%
##
## County data found: 100 unique counties
##
## Creating initial partition...
## Anti-Bunking List 1: weight=1000.0, exponent=1.0, 14 precincts
## GUIDED cooling: T0=409.4 to T=1.0 at iteration 2250 (90%) at rate=0.997330
## Optimization enabled: Initial score=2047.0, Temp=409.4
## County-aware mode: bias factor = 3.0
##
## Running 2500 ReCom steps...
## Iter: 100/2500
## Iter: 200/2500
## Iter: 300/2500
## Iter: 400/2500
```

```

## Iter: 500/2500
## Iter: 600/2500
## Iter: 700/2500
## Iter: 800/2500
## Iter: 900/2500
## Iter: 1000/2500
## Iter: 1100/2500
## Iter: 1200/2500
## Iter: 1300/2500
## Iter: 1400/2500
## Iter: 1500/2500
## Iter: 1600/2500
## Iter: 1700/2500
## Iter: 1800/2500
## Iter: 1900/2500
## Iter: 2000/2500
## Iter: 2100/2500
## Iter: 2200/2500
## Iter: 2300/2500
## Iter: 2400/2500
## Iter: 2500/2500
##
##
## Completed in 46.5s (18.6ms/step, 53.7 iter/s)
## Optimization stats: 1104 accepts, 1396 rejects (44.2% accept rate)
## Score: 2047.0 to 828.9 (D=-1218.2)
##
## =====
## FINAL ANALYSIS
## =====
## District      DEM       REP      DEM%      REP%      Margin Winner
## -----
## D1          149158    169646   46.8%     53.2%     6.4% REP
## D2          248386    130408   65.6%     34.4%    31.1% DEM
## D3          203585    206033   49.7%     50.3%     0.6% REP
## D4          175137    264429   39.8%     60.2%    20.3% REP
## D5          255717    179687   58.7%     41.3%    17.5% DEM
## D6          147139    241142   37.9%     62.1%    24.2% REP
## D7          260045    152935   63.0%     37.0%    25.9% DEM
## D8          222738    173636   56.2%     43.8%    12.4% DEM
## D9          123609    288234   30.0%     70.0%    40.0% REP
## D10         200942    246447   44.9%     55.1%    10.2% REP
## D11         202911    202496   50.1%     49.9%     0.1% DEM
## D12         162791    240380   40.4%     59.6%    19.2% REP
## D13         200248    168350   54.3%     45.7%     8.7% DEM
## D14         161203    233118   40.9%     59.1%    18.2% REP
## -----
## TOTAL        2713609   2896941  48.4%     51.6%
##
## Mean Dem Share: 48.4%
## Median Dem Share: 48.2%
## MM: R+0.2% (tgt R+0.0%) | MM Score: 0
## EG: R+3.0% (tgt R+0.0%) | EG Score: 0
##
## County Splits Score: 9.40
##
## -----
## Anti-Bunking Metrics:
## List 1: 14 precincts in 13 districts | metric=0.148 | score=148.4 (w=1000 e=1.0)

```

```

## =====
## 
## =====
## RESULTS SAVED
## =====
## Final assignment: output/final_assignment_12182025_0039.csv
## Metrics: output/metrics_12182025_0039.csv
## Snapshots: output/snapshot_assignments_12182025_0039.csv
## =====

```

Partisanship looks good. No major mean/median skew, and a healthy number of competitive districts.

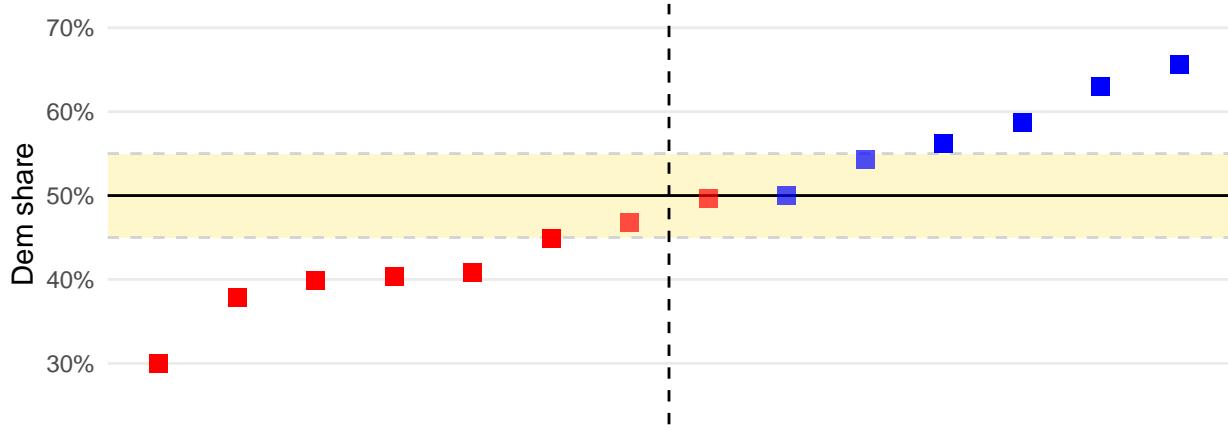
```

## Loading data...
## Using most recent metrics CSV: output/metrics_12182025_0039.csv
## Loading metrics for partisan plot...
## Rendering partisan graph plot...

```

## Example 6

### Fair NC Congressional Map



We can also pass bunking lists to `mosaic_plot()` - which makes it easy to visually check out the results.

```

p <- mosaic_plot("shapefiles/North_Carolina_Simplified.shp",
                  bunking_lists = create_bunking_lists(
                    anti_bunking = list(NC_incumbent_list)
                  ),
                  title = "Example 6: Fair NC Congressional Map Without Double-Bunking",
                  subtitle = "Yellow diamonds: incumbent locations",
                  district_outline = TRUE,
                  number_labels = TRUE,
                  county_outline = TRUE,
                  border_outline = TRUE)
p

```

## Example 6: Fair NC Congressional Map Without Double-Bunking

Yellow diamonds: incumbent locations

